

```
1 !pip install -Uq transformers accelerate einops bitsandbytes chromadb langchain pymupdf pytz gradio
```

```
📦 _____ 8.8/8.8 MB 25.9 MB/s eta 0:00:00
_____ 297.4/297.4 kB 12.4 MB/s eta 0:00:00
_____ 44.6/44.6 kB 3.0 MB/s eta 0:00:00
_____ 102.2/102.2 MB 5.9 MB/s eta 0:00:00
_____ 525.5/525.5 kB 18.5 MB/s eta 0:00:00
_____ 814.5/814.5 kB 7.8 MB/s eta 0:00:00
_____ 3.9/3.9 MB 18.0 MB/s eta 0:00:00
_____ 505.5/505.5 kB 16.9 MB/s eta 0:00:00
_____ 17.1/17.1 MB 33.1 MB/s eta 0:00:00
_____ 2.4/2.4 MB 24.4 MB/s eta 0:00:00
_____ 91.9/91.9 kB 5.0 MB/s eta 0:00:00
_____ 60.8/60.8 kB 5.7 MB/s eta 0:00:00
_____ 41.3/41.3 kB 3.1 MB/s eta 0:00:00
_____ 5.4/5.4 MB 27.3 MB/s eta 0:00:00
_____ 6.8/6.8 MB 41.7 MB/s eta 0:00:00
_____ 60.1/60.1 kB 3.8 MB/s eta 0:00:00
_____ 106.1/106.1 kB 9.1 MB/s eta 0:00:00
_____ 67.3/67.3 kB 4.6 MB/s eta 0:00:00

Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing metadata (pyproject.toml) ... done
_____ 698.9/698.9 kB 44.2 MB/s eta 0:00:00
_____ 1.6/1.6 MB 49.7 MB/s eta 0:00:00
_____ 67.6/67.6 kB 8.4 MB/s eta 0:00:00
_____ 144.8/144.8 kB 18.9 MB/s eta 0:00:00
_____ 1.9/1.9 MB 83.8 MB/s eta 0:00:00
_____ 278.4/278.4 kB 34.1 MB/s eta 0:00:00
_____ 104.2/104.2 kB 13.9 MB/s eta 0:00:00
_____ 30.8/30.8 MB 36.1 MB/s eta 0:00:00

Preparing metadata (setup.py) ... done
_____ 313.6/313.6 kB 34.1 MB/s eta 0:00:00
_____ 75.6/75.6 kB 10.4 MB/s eta 0:00:00
_____ 8.7/8.7 MB 54.0 MB/s eta 0:00:00
_____ 129.9/129.9 kB 16.6 MB/s eta 0:00:00
_____ 49.4/49.4 kB 5.8 MB/s eta 0:00:00
_____ 71.9/71.9 kB 9.1 MB/s eta 0:00:00
_____ 77.9/77.9 kB 10.0 MB/s eta 0:00:00
_____ 58.3/58.3 kB 7.8 MB/s eta 0:00:00
_____ 53.0/53.0 kB 6.8 MB/s eta 0:00:00
_____ 46.0/46.0 kB 6.4 MB/s eta 0:00:00
_____ 50.8/50.8 kB 7.1 MB/s eta 0:00:00
_____ 341.4/341.4 kB 38.5 MB/s eta 0:00:00
_____ 3.4/3.4 MB 61.8 MB/s eta 0:00:00
_____ 1.3/1.3 MB 80.2 MB/s eta 0:00:00
_____ 86.8/86.8 kB 11.9 MB/s eta 0:00:00

Building wheel for pypika (pyproject.toml) ... done
Building wheel for ffmpeg (setup.py) ... done
```

```
1 import json
2 import gradio
3
4 import transformers
5 from transformers import AutoModelForCausalLM, AutoTokenizer, pipeline
6
7 import torch
8 from torch import cuda, bfloat16
9 from langchain.text_splitter import RecursiveCharacterTextSplitter
10
11 import chromadb
12 from chromadb.utils import embedding_functions
13
14 from datetime import datetime
15 import json
16 import fitz
17 import pytz
18
19 import os
20
21 # if using Google Colab uncomment the following lines
22 # from google.colab import drive
23
24 # drive.mount("/content/drive")
25 # os.chdir(r"/content/drive/My Drive/AAI courses/TA Chatbot") # Change to desired dir
```

Mounted at /content/drive

Set up the database

```
1 CHROMA_DATA_PATH = "chromadb_database"
2
3
4 def fetch_database(hf_auth=None, ef_name="sentence-transformers/all-mpnet-base-v2"):
5     if hf_auth:
6         ef = embedding_functions.HuggingFaceEmbeddingFunction(
7             api_key=hf_auth, model_name=ef_name
8         )
9     else:
10        ef = (
11            embedding_functions.DefaultEmbeddingFunction()
12        ) # use the default all-MiniLM-L6-v2
13
14    client = chromadb.PersistentClient(path=CHROMA_DATA_PATH)
15    collection = client.get_or_create_collection(
16        name="course_material", embedding_function=ef
17    )
18    return client, collection
19
20
21 # with open("creds.txt", "r") as file: # load your hugging face token
22 #     hf_auth = json.load(file)["hf_token"]
23
24 client, collection = fetch_database()
```

```
1 def load_material(collection=collection, chunk_size=250, chunk_overlap=20):
2     text_splitter = RecursiveCharacterTextSplitter(
3         chunk_size=chunk_size, chunk_overlap=chunk_overlap
4     )
5     if collection.count(): # if the store is already built, abort
6         return
7     ids = 0
8     data_path = "data"
9
10    for file in os.listdir(data_path):
11        docs = []
12        if file.startswith(
13            "module"
14        ): # expecting data poisoning folders to have the name module#
15            week_num = file[6]
16        else:
17            week_num = 0 # means available at all time
18
19        try:
20            curr_file = fitz.open(os.path.join(data_path, file))
21            for i in range(len(curr_file)):
22                if curr_file[i].get_text():
23                    for doc in text_splitter.split_text(curr_file[i].get_text()):
24                        docs.append(doc)
25            idx = list(f"{i}" for i in range(ids, ids + len(docs)))
26            ids += len(docs)
27            meta_data = [{"week": f"{week_num}"}] * len(idx)
28
29            collection.add(documents=docs, metadatas=meta_data, ids=idx)
30
31    except Exception as e:
32        print("Could not process ", file, "\nError: ", e, sep="")
33
34
35 load_material()
```

```

1 # set the end date for each week
2 dates = [
3     datetime.strptime("2024-04-09", "%Y-%m-%d").date(), # beginning of week 1
4     datetime.strptime("2024-04-10", "%Y-%m-%d").date(), # end of week 1
5     datetime.strptime("2024-04-11", "%Y-%m-%d").date(), # end of week 2
6     datetime.strptime("2024-04-12", "%Y-%m-%d").date(), # end of week 3...
7     datetime.strptime("2024-04-13", "%Y-%m-%d").date(),
8     datetime.strptime("2024-04-14", "%Y-%m-%d").date(),
9     datetime.strptime("2024-04-15", "%Y-%m-%d").date(),
10 ]
11
12
13 def query_db(
14     query, metadata_filters, n_results=1
15 ): # only fetch the most similar document
16     results = collection.query(
17         query_texts=query,
18         n_results=n_results,
19         where=metadata_filters,
20     )
21     return results
22
23
24 tz = pytz.timezone("America/Los_Angeles")
25
26
27 def fetch_docs(query, n_docs=3, tz=tz):
28     now = datetime.now(tz).date()
29
30     for i in range(len(dates)):
31         if now <= dates[i]:
32             metadata_filters = {"week": {"$in": ["0", str(i + 1)]}}
33
34             return "\n\n".join(
35                 query_db(query, metadata_filters, n_docs)["documents"][0]
36             )

```

Load the LLM

```

1 def load_model(hf_auth, model_id, temperature=0.001):
2     # Determine the device (GPU if available, else CPU)
3     device = f"cuda:{cuda.current_device()}" if cuda.is_available() else "cpu"
4
5     # Configure quantization settings for loading the model with less GPU memory usage
6     bnb_config = transformers.BitsAndBytesConfig(
7         load_in_4bit=True,
8         bnb_4bit_quant_type="nf4",
9         bnb_4bit_use_double_quant=True,
10        bnb_4bit_compute_dtype=bfloat16,
11    )
12
13    # Load the configuration for the pre-trained model
14    model_config = transformers.AutoConfig.from_pretrained(model_id, token=hf_auth)
15
16    tokenizer = AutoTokenizer.from_pretrained(model_id, token=hf_auth)
17
18    # Load the model for causal language modeling
19    model = transformers.AutoModelForCausalLM.from_pretrained(
20        model_id,
21        trust_remote_code=True,
22        config=model_config,
23        quantization_config=bnb_config,
24        device_map="auto",
25        token=hf_auth,
26    )
27
28    # Set the model in evaluation mode for inference
29    model.eval()
30
31    pipe = pipeline(
32        task="text-generation",
33        model=model,
34        tokenizer=tokenizer,
35        temperature=temperature,
36        # top_k=40,
37        eos_token_id=tokenizer.eos_token_id,
38        # repetition_penalty=1.5,
39        return_full_text=False,
40        # do_sample=True,
41    )
42
43    # Print device information where the model is loaded
44    print(f"Model loaded on {device}")
45
46    return pipe
47
48
49 with open("creds.txt", "r") as file:
50     hf_auth = json.load(file)["hf_token"]
51
52 pipe = load_model(hf_auth, "meta-llama/Llama-2-13b-chat-hf")

```

Loading checkpoint shards: 100%

2/2 [04:30<00:00, 121.50s/it]

Model loaded on cuda:0

We will use the following prompt template:

```

<s>[INST] <<SYS>>
{{ system_prompt }}
</SYS>>

{{ user_msg_1 }} [/INST] {{ model_answer_1 }} </s><s>[INST] {{ user_msg_2 }} [/INST] {{ model_answer_2 }} </s><s>[INST] {{ user_msg_3 }} [/INST]

```

We will pass the context in the following structure:

Context information from multiple sources is below.

{context_str}

Given the information from multiple sources and not prior knowledge, answer the query.

Query: {query_str}

Answer:

```
1 SYS_PROMPT = """"You are an expert teacher assistant for a course called ADS 500B. You answer questions relating to software engineering,
2 You are honest and helpful. You answer succinctly and professionally. You do not make up facts. Use the Context provided to answer the qu
3 If you're asked something you do not know the answer to, say you do not know. Do not make up facts. Be brief and to the point.\"
4 ""\"
5
6 TEMPLATE = "<s>[INST] <<SYS>>\n{sys_prompt}\n<</SYS>>\"
7
8 _SYS_PROMPT = TEMPLATE.format(sys_prompt=SYS_PROMPT)
9
10
11 def build_prompt(query_w_context: str, hist_len: int = 3):
12     if not len(history): # if first query
13         return _SYS_PROMPT + f\"\\n\\n{query_w_context} [/INST] \"
14
15     prompt = _SYS_PROMPT + \"\\n\\n\"
16
17     for i in range(len(history)):
18         if i == 0:
19             prompt += f\"{history[i]} [/INST] \"
20             continue
21
22         if i % 2 == 0:
23             prompt += f\"<s>[INST]{history[i]} [/INST] \"
24
25         else:
26             prompt += f\"{history[i]} </s>\"
27
28     return prompt + f\"<s>[INST]{query_w_context} [/INST] \"
29
30
31 def query_model(query: str, hist_len: int = 3):
32     global history
33
34     context = fetch_docs(query=query, n_docs=1)
35
36     query_w_context = f\"\"\"Context information from multiple sources is below.
37 -----
38 {context}
39 -----
40 Given the information from multiple sources and not prior knowledge, answer the query.
41 Query: {query}
42 Answer: \"\"\"
43
44     prompt = build_prompt(query_w_context, hist_len)
45     response = pipe(prompt, max_new_tokens=1024)[0][\"generated_text\"]
46
47     history.append(query)
48     history.append(response)
49
50     if len(history) > hist_len:
51         history = history[
52             -(hist_len * 2) :
53             ] # only keep hist_len interaction pair history.
54     return response
```

```
1 global history
2 history = []
3 hist_len = 2
4
5
6 def chatbot_interface(query, gradio_hist):
7     gradio_hist = None
8     result = query_model(query)
9     return result
10
11
1 interface.launch(share=True, debug=True)
```

Colab notebook detected. This cell will run indefinitely so that you can see errors and logs. To turn off, set debug=False in launch().
Running on public URL: <https://3a92f59170c09a1ddd.gradio.live>

This share link expires in 72 hours. For free permanent hosting and GPU upgrades, run `gradio deploy` from Terminal to deploy to Spaces